

Starting points for future work

L Gatto and S J Eglén

October 29, 2012

The shell

The unix shell is an extremely powerful environment that features many extremely handy tools, that do simple things, and that can be piped (|) together.

- ▶ wc, grep, cut
- ▶ tr, sed, awk

shell can also be used for scripting.

Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

grep, sed, awk

- ▶ All exploit regular expressions. See ItDT book (later).
- ▶ grep: find matching lines
- ▶ sed: stream-editor. Incredibly handy for one-liners:

<http://sed.sourceforge.net/sed1line.txt>

```
sed 's/foo/bar/g'           # replaces ALL instances in line
# print section of file between two regular expressions
sed -n '/Iowa/,/Montana/p'  # case sensitive
```

- ▶ awk: flexible pattern matching/ processing of text files.

<http://www.pement.org/awk/awk1line.txt>

```
# print the sums of the fields of every line
awk '{s=0; for (i=1; i<=NF; i++) s=s+$i; print s}'
```

diff: where do my files differ?

version1.dat

```
0.701 -0.764 -0.226 0.796 -0.337
0.249 -1.51 0.876 2.25 -0.879
-0.523 -1.29 0.354 -0.378 -1.39
0.565 1.31 -0.237 -0.844 0.28
2 -0.128 -0.841 1.31 -0.651
-0.565 0.81 -0.116 0.582 -0.0334
1.03 -0.75 1.7 -0.829 2.3
0.797 -0.988 0.667 -0.492 -0.78
0.94 -0.0931 -0.22 -1.29 -1.21
-0.456 -0.0231 0.603 1.43 0.734
0.598 -0.113 0.852 -1.58 -0.165
0.126 -0.0806 0.951 0.49 0.328
```

version2.dat

```
0.701 -0.764 -0.226 0.796 -0.337
0.249 -1.51 0.876 2.25 -0.879
-0.523 -1.29 0.354 -0.378 -1.39
0.565 1.31 -0.235 -0.844 0.28
2 -0.128 -0.841 1.31 -0.651
-0.565 0.81 -0.116 0.582 -0.0334
1.03 -0.75 1.7 -0.829 2.3
0.797 -0.988 0.667 -0.492 -0.78
0.94 -0.0932 -0.22 -1.29 -1.21
-0.456 -0.0231 0.603 1.43 0.734
0.598 -0.113 0.852 -1.58 -0.165
0.126 -0.0806 0.951 0.49 0.328
```

diff and patch

- ▶ *diff* shows the differences between version1 and version 2.

```
diff nextsteps/version1.dat nextsteps/version2.dat
```

- ▶ *patch*: new file = old file + diff
- ▶ *patches* are efficient ways of sending updates. Useful for syncing and version control.

```
diff version1.dat version2.dat > p
patch version1.dat p
diff version1.dat version2.dat
```

Perl: Practical Extraction and Report Language

- ▶ Most unix tools (used to be) limited by length of lines. Perl removed those restrictions, combining features of awk, sh and C.
- ▶ ‘duct tape’ programming language.
- ▶ Useful in computational biology. See <http://www.bioperl.org>
- ▶ Excellent Ensembl API, <http://www.ensembl.org/info/data/api.html>
- ▶ G. Valiente. Combinatorial Pattern Matching Algorithms in Computational Biology using Perl and R . Taylor & Francis/CRC Press (2009).
- ▶ Verdict: yucky, but probably [essential | good to now].
- ▶ Bidirectional R /Perl interfaces <http://www.omegahat.org/RSPerl/>

R can also regexp

- ▶ `grep`, `sub`, `gsub`, `strsplit`, `nchar`, `substr`, ...
- ▶ also `stringr` package

and for sequence data storing and manipulation

- ▶ `Biostrings` package

Python

- ▶ Modern programming language; less compact than perl:

```
while (<>) {           | import sys
    print if /perl/i;   | for line in sys.stdin.readlines():
}                       |     if line.lower().find("perl") > -1:
                        |         print line,
```

<http://www.sabren.net/articles/againstperl.php3>

- ▶ Clean syntax
- ▶ Properly object-oriented.
- ▶ Not as much support in computational biology (yet). See <http://www.biopython.org>
- ▶ Verdict: More general programming language than R ; lacking (perhaps?) in core numerics and graphics – see NumPy and RPy(2).
- ▶ Bidirectional R /Python interface <http://www.omegahat.org/RSPython/>

C

- ▶ Low-level programming language
- ▶ Very fast, but takes a long time to write code.
- ▶ You have to worry about memory allocation yourself.
- ▶ All variables have predefined type.
- ▶ Critical for numerical-intensive work. (FORTRAN less-popular.)

Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

C from R

- ▶ R has build-in C interfaces
 - ▶ Better know how to program in C.
 - ▶ Documentation is not always easy to follow: R-Ext, R Internals as well as R and other package's code.
- ▶ .C Arguments and return values must be *primitive* (vectors of doubles or integers)
- ▶ .Call Accepts R data structures as arguments and return values (SEXP and friends) (no type checking is done though).
- ▶ Memory management: memory allocated for R objects is garbage collected. Thus R objects in C code, you must be explicitly PROTECTed to avoid being gc()ed, and subsequently UNPROTECTed.

Using .Call

```
#include <R.h>
#include <Rdefines.h>

SEXP gccount(SEXP inseq) {
  int i, l;
  SEXP ans, dnaseq;
  PROTECT(dnaseq = STRING_ELT(inseq, 0));
  l = LENGTH(dnaseq);
  printf("length %d\n", l);
  PROTECT(ans = NEW_NUMERIC(4));

  for (i = 0; i < 4; i++)
    REAL(ans)[i] = 0;

  for (i = 0; i < l; i++) {
    char p = CHAR(dnaseq)[i];
    if (p=='A')
      REAL(ans)[0]++;
    else if (p=='C')
      REAL(ans)[1]++;
    else if (p=='G')
      REAL(ans)[2]++;
    else if (p=='T')
      REAL(ans)[3]++;
    else
      error("Wrong alphabet");
  }
  UNPROTECT(2);
  return(ans);
}
```

Rcpp for C++

- ▶ Rcpp is a great package for writing both C and C++ code:
- ▶ It comes with loads of documentation and examples.
- ▶ No need to worry about garbage collection.
- ▶ All basic R types are implemented as C++ classes.
- ▶ Easy to interface C++ classes (via modules)
- ▶ With package inline code can be easily compiled in R .

```
library(Rcpp)
library(inline)
cppCode <- '
  Rcpp::NumericVector cx(x);
  Rcpp::NumericVector ret(1);
  ret[0] = cx[0] * cx[0];
  return(ret);
',

squareOne <- cxxfunction(signature(x="numeric"),
                          plugin="Rcpp", body=cppCode)

squareOne(10)
```

Using our C code

- ▶ Create a shared library: R CMD SHLIB gccount.c
- ▶ Load the shared object: dyn.load("gccount.so")
- ▶ Create an R function that uses it: gccount <- function(inseq) .Call("gccount",inseq)
- ▶ Use the C code: gccount("GACAGCATCA")

```
s <- "GACTACGA"
gccount
gccount(s)
table(strsplit(s, ""))
system.time(replicate(10000, gccount(s)))
system.time(replicate(10000, table(strsplit(s, ""))))
```

Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

Syncing your files

- ▶ How do you keep two directories in synchrony, e.g. your home directory on laptop and desktop?
- ▶ `sftp`, `ssh`, `rsync`
- ▶ Unison gets Stephen's vote since 2003 –
<http://www.damtp.cam.ac.uk/internal/computing/unison/>
- ▶ Modern services like Dropbox are useful and build upon these unix tools.

For packages, analysis projects, papers and slides

- ▶ Github, google code, bitbucket, ...
- ▶ R-forge: `svn` and build system

Version control / revision control system (RCS)

- ▶ How to keep backup copies over time?
- ▶ Just copy files, e.g. *mycode.jan1.R*, *mycode.jan2.R*, ...
- ▶ Leads to many large copies, with no trace of what you did over time.
- ▶ more principled way is to use version control: every time you make significant changes, you *commit* a new version with a succinct log file saying what you changed.
- ▶ RCS: going since 1982... old and simple but stable. Typically single-user.

<http://www.cl.cam.ac.uk/~mgk25/rcsintro.html>

- ▶ More modern approaches: *cvs*, *svn*, *git*, ...

Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

Handling large data files.

- ▶ Computational Biology requires access to large data files.
- ▶ Reading them all into memory is difficult, when files are very large (> 1 Gb).
- ▶ Some approaches:
 1. Compress files.
 2. Selectively use scan or connections.
 3. Use a database.

2. Scan and Connections.

- ▶ `scan()` is very flexible; e.g. read just 2nd column:

```
scan(file = "", what = double(0), nmax = -1, n = -1, sep = "",
      quote = if(identical(sep, "\n")) "" else "'", dec = ".",
      skip = 0, nlines = 0, na.strings = "NA",
      flush = FALSE, fill = FALSE, strip.white = FALSE,
      quiet = FALSE, blank.lines.skip = TRUE, multi.line = TRUE,
      comment.char = "", allowEscapes = FALSE,
      fileEncoding = "", encoding = "unknown")
```

```
x <- scan(file, what=list(NULL,"",NULL), skip=2, sep='\t')
```

- ▶ connections allow you to maintain state between accesses to a file.

```
con <- file("version1.dat", "r")
while (length(dat <- scan(con,n = 5,quiet = TRUE)) > 0) {
  print(mean(dat))
}
close(con)
```

1. Compress files.

- ▶ This produces typically x2 compression:

```
Rscript -e 'write(rnorm(99999), file="largefile.dat")'
ls -lh largefile.dat
gzip largefile.dat
ls -lh largefile.dat.gz
gunzip largefile.dat
```

- ▶ R can read in compressed files natively.

```
x <- scan('largefile.dat.gz')
```

- ▶ Other compression options also recognised: xz, bzip2

3. Relational databases

- ▶ Relational database: data stored in tables, very similar in nature to R's data.frames.
- ▶ Databases allow for multiple-accesses, locks for restricted changes, very scalable.
- ▶ Many databases available: Oracle, Postgres, Access, MySQL.
- ▶ SQL – Structured Query Language: language to interrogate databases.

What is SQLite?

- ▶ Most databases run on remote server; SQLite is embedded into your program.
- ▶ Embedding the database simplifies setup of server, but means your databases are not shared in the same way that others are. (You have to share the .sql files.)
- ▶ Incredibly small (1/4 Mb) and useful. Widely used (e.g. mac, iOS, Firefox, Android). Not as fast as e.g. Oracle.
- ▶ You compile your SQLite within your program.
- ▶ All handled with you by R, care of *RSQLite* package. (e.g. Bioconductor uses it for data files.)

Other uses for sqlite

- ▶ sqldf Performs SQL selects on R data frames.
- ▶ supports SQLite backend database (by default), the H2 java db and PostgreSQL and MySQL.
- ▶ avoid read.csv entirely <http://code.google.com/p/sqldf/>

"See ?read.csv.sql in sqldf. It uses RSQLite and SQLite to read the file into an sqlite database (which it sets up for you) completely bypassing R and from there grabs it into R removing the database it created at the end." (G. Grothendieck, r-help mailing list).

- ▶ Good book: $\sim((HT|X)M|SQ)L|R\$$
Introduction to Data Technologies.

<http://www.stat.auckland.ac.nz/~paul/ItDT/>

Using databases in R, a simple session (Gentleman, p239)

- ▶ package *DBI* interfaces to all database platforms.

```
library(RSQLite)
m = dbDriver("SQLite")

## Create a new database from an R data frame.
con = dbConnect(m, dbname = "arrest.db")
data(USArrests)
dbWriteTable(con, "USArrests", USArrests, overwrite=TRUE)
dbListTables(con)

## Later, query the database.
rs = dbSendQuery(con, "select * from USArrests")
d1 = fetch(rs, n=5)      ## get first five
print(d1)
d1 = fetch(rs, n=-1)
dbDisconnect(con)
```

ff: back to the future?

- ▶ *ff* package stores objects on disk, but looks like they are in memory.
- ▶ "back to the future": S used to store objects in disk.
- ▶ Sorting a single column of 81e6 entries. Time-taken in seconds.

Oct 2010 results from.

<http://tolstoy.newcastle.edu.au/R/packages/10/0697.html>

	ruinteger	rinteger	rusingle	rsingle	rudouble	rdouble	rfactor
ram	5.58	3.23	NA	NA	NA	NA	0.49
ff	10.70	8.54	51.35	28.98	70.20	44.13	7.91
R	OOM	OOM	OOM	OOM	OOM	OOM	OOM
SAS	61.45	44.94	NA	NA	63.14	46.56	NA

(ram=in-memory, optimized for speed, not ram; ff=on disk).

Other

- ▶ The bigmemory package by Kane and Emerson permits storing large objects such as matrices in memory (as well as via files) and uses external pointer objects to refer to them.
- ▶ netCDF data files: ncdcf and RNetCDF packages.
- ▶ hdf5 format: rhdf5 package

Introduction

- ▶ Applicable when repeating *independent* computations a certain number of times; results just need to be combined after parallel executions are done.
- ▶ A cluster of nodes: generate multiple workers listening to the master; these workers are new processes that can run on the current machine or a similar one with an identical R installation. Should work on all R platforms (as in package snow).
- ▶ The R process is *forked* to create new R ~processes by taking a complete copy of the masters process, including workspace (pioneered by package multicore). Does not work on Windows.
- ▶ Grid computing.

Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

Packages

- ▶ Package parallel, first included in R 2.14.0 builds on CRAN packages multicore and snow.

```
mclapply(X, FUN, ...) (adapted from multicore).  
parLapply(cl, X, FUN, ...) (adapted from snow ).
```

- ▶ Package foreach, introducing a new looping construct supporting parallel execution. Natural choice to parallelise a for loop.

```
library(doMC)  
library(foreach)  
registerDoMC(2)  
foreach(i = 1:10) %dopar% f(i)  
foreach(i = 1:10) %do% f(i) ## serial version  
library(plyr)  
lply(1:10, f, .parallel=TRUE)
```


Outline

Key unix tools and languages

C and C++ from R

Not only local

Handling large files / databases

Parallel processing

A few more words about about R

Conclusions

Profiling

```
m <- matrix(rnorm(1e6), ncol=100)
Rprof("rprof")
res <- apply(m,1,mean,trim=.3)
Rprof(NULL)
summaryRprof("rprof")
```

Pass by ...

- ▶ **value** is the default in R
- ▶ **reference** using S4 ReferenceClasses (OO)
- ▶ can emulate pass by ref using an environment

```
e <- new.env()
e$x <- 1
f <- function(myenv) myenv$x <- 2
f(e)
e$x
```

Benchmarking

```
m <- matrix(rnorm(1e6), ncol=100)
f1 <- function(x, t = 0.3) {
  xx <- 0
  for (i in 1:nrow(x)) {
    xx <- c(xx, sum(m[i, ]))
  }
  mean(xx, trim = t)
}
f2 <- function(x, t = 0.3) mean(rowSums(x), trim = t)

library(rbenchmark)
benchmark(f1(m), f2(m),
          columns=c("test", "replications",
                    "elapsed", "relative"),
          order = "relative", replications = 10)
```