

Data structures

2

R is designed to handle experimental data

- Although the basic unit of R is a vector, we usually handle data in **data frames**.
- A data frame is a set of observations of a set of variables – in other words, the outcome of an experiment.
- For example, we might want to analyse information about a set of patients. To start with, let's say we have ten patients and for each one we know their name, sex, age, weight and whether they give consent for their data to be made public.
- Load this data into a data frame called 'patients' in R:

```
source("05_patients.R")
```

The patients data frame

- The 'patients' data frame has ten rows (observations) and seven columns (variables). The columns must all be equal lengths.

```
> patients
```

	First_Name	Second_Name	Full_Name	Sex	Age	Weight	Consent
1	Adam	Jones	Adam Jones	Male	50	70.8	TRUE
2	Eve	Parker	Eve Parker	Female	21	67.9	TRUE
3	John	Evans	John Evans	Male	35	75.3	FALSE
4	Mary	Davis	Mary Davis	Female	45	61.9	TRUE
5	Peter	Baker	Peter Baker	Male	28	72.4	FALSE
6	Paul	Daniels	Paul Daniels	Male	31	69.9	FALSE
7	Joanna	Edwards	Joanna Edwards	Female	42	63.5	FALSE
8	Matthew	Smith	Matthew Smith	Male	33	71.5	TRUE
9	David	Roberts	David Roberts	Male	57	73.2	FALSE
10	Sally	Wilson	Sally Wilson	Female	62	64.8	TRUE

- Let's see how we can construct this from scratch.

Character, numeric and logical data types

- Each column is a vector, like previous vectors we have seen:

```
> firstName<- c("Adam", "Eve", "John", "Mary", "Peter", "Paul", "Joanna",  
"Matthew", "David", "Sally")  
> secondName<-c("Jones", "Parker", "Evans", "Davis", "Baker", "Daniels",  
"Edwards", "Smith", "Roberts", "Wilson")  
> age<-c(50, 21, 35, 45, 28, 31, 42, 33, 57, 62)  
> weight<-c(70.8, 67.9, 75.3, 61.9, 72.4, 69.9, 63.5, 71.5, 73.2, 64.8)  
> consent<-c(TRUE,TRUE,FALSE,TRUE,FALSE,FALSE,FALSE,TRUE,FALSE,TRUE)
```

- Each vector has a type, which we can see with the **mode** function:

```
> mode(firstName)  
[ 1] "character"  
> mode(age)  
[ 1] "numeric"  
> mode(weight)  
[ 1] "numeric"  
> mode(consent)  
[ 1] "logical"
```

Factors

- Character vectors are fine for some variables, like names
- But sometimes we have categorical data and we want R to recognise this.
- A factor is R's data structure for categorical data.

```
> sex<-c("Male", "Female", "Male", "Female", "Male", "Male", "Female",  
"Male", "Male", "Female")  
> sex  
[1] "Male" "Female" "Male" "Female" "Male" "Male" "Female" "Male"  
"Male" "Female"  
> factor(sex)  
[1] Male Female Male Female Male Male Female Male Male Female  
Levels: Female Male
```

- R has converted the strings of the sex character vector into two **levels**, which are the categories in the data.
- Note the values of this factor are not character strings, but levels.
- We can use this factor to compare data for males and females.

Creating a data frame (first attempt)

- We can construct a data frame from other objects:

```
> patients<-data.frame(firstName, secondName, paste(firstName,secondName),  
sex, age, weight, consent)  
> patients  
  firstName secondName paste.firstName..secondName. sex age weight consent  
1      Adam      Jones      Adam Jones      Male  50   70.8    TRUE  
2        Eve      Parker      Eve Parker Female  21   67.9    TRUE  
3       John      Evans      John Evans      Male  35   75.3    FALSE  
4        Mary      Davis      Mary Davis Female  45   61.9    TRUE  
5       Peter      Baker      Peter Baker      Male  28   72.4    FALSE  
6        Paul      Daniels      Paul Daniels      Male  31   69.9    FALSE  
7      Joanna      Edwards      Joanna Edwards Female  42   63.5    FALSE  
8     Matthew      Smith      Matthew Smith      Male  33   71.5    TRUE  
9       David      Roberts      David Roberts      Male  57   73.2    FALSE  
10      Sally      Wilson      Sally Wilson Female  62   64.8    TRUE
```

- The **paste** function joins character vectors together.
- We can access particular variables using the **dollar** operator:

```
> patients$age  
[1] 50 21 35 45 28 31 42 33 57 62
```

Naming data frame variables

- R has inferred the names of our data frame variables from the names of the vectors or the commands (eg the paste command).
- We can name the variables after we have created a data frame using the **names** function, and we can use the same function to see the names:

```
> names(patients)<-c("First_Name", "Second_Name", "Full_Name", "Sex",  
"Age", "Weight", "Consent")  
> names(patients)  
[1] "First_Name" "Second_Name" "Full_Name" "Sex" "Age"  
"Weight" "Consent"
```
- Or we can name the variables when we define the data frame:

```
> patients<-data.frame(First_Name=firstName, Second_Name=secondName,  
Full_Name=paste(firstName,secondName), Sex=sex, Age=age, Weight=weight,  
Consent=consent)  
> names(patients)  
[1] "First_Name" "Second_Name" "Full_Name" "Sex" "Age"  
"Weight" "Consent"
```

Factors in data frames

- When creating a data frame, R assumes all character vectors should be categorical variables and converts them to factors. This is not always what we want:

```
> patients$firstName  
[1] Adam Eve John Mary Peter Paul Joanna Matthew David Sally  
Levels: Adam David Eve Joanna John Mary Matthew Paul Peter Sally
```
- We can avoid this by asking R not to treat strings as factors, and then explicitly stating when we want a factor by using **factor**:

```
> patients<-data.frame(First_Name=firstName, Second_Name=secondName,  
Full_Name=paste(firstName,secondName), Sex=factor(sex), Age=age,  
Weight=weight, Consent=consent, stringsAsFactors=FALSE)  
> patients$Sex  
[1] Male Female Male Female Male Male Female Male Male Female  
Levels: Female Male  
> patients$firstName  
[1] "Adam" "Eve" "John" "Mary" "Peter" "Paul" "Joanna"  
"Matthew" "David" "Sally"
```

Storage modes & data types

- Data types - why care?
 - May get an undesired result if calculations are between numbers stored as different types
- R will coerce data types when calculations between differing types are forced
 - If the operation is inappropriate, the calculation will fail.
e.g.

```
> 2 + "2"
```


will fail as we cannot add a character string to integer!

Matrices

`matrix(..., ncol=..., nrow=...)`

- Data frames are R's speciality, but R also handles matrices:

```
> e <- matrix(1:10, nrow=5, ncol=2)
> e
     [,1] [,2]
[1,]    1    6
[2,]    2    7
[3,]    3    8
[4,]    4    9
[5,]    5   10
> f <- matrix(1:10, nrow=2, ncol=5)
> f
     [,1] [,2] [,3] [,4] [,5]
[1,]    1    3    5    7    9
[2,]    2    4    6    8   10
> f %*% e
     [,1] [,2]
[1,]   95   220
[2,]  110   260
```

The `%*%` operator is the matrix multiplication operator, not the standard multiplication operator.

Indexing data frames and matrices

Special cases:
`a[i,]` i-th row
`a[,j]` j-th column

- You can index multidimensional data structures like matrices and data frames using commas. If you don't provide an index for either rows or columns, all of the rows or columns will be returned.

`object [rows , columns]`

```
> e[1,2]
[1] 6
> e[1,]
[1] 1 6
> patients[1,2]
[1] "Jones"
> patients[1,]
  First_Name Second_Name Full_Name Sex Age Weight Consent
1      Adam      Jones Adam Jones Male  50   70.8    TRUE
```

Advanced indexing

- As values in R are really vectors, so indices are actually vectors, and can be numeric or logical:

```
> s <- letters[1:5]
> s[c(1,3)]
[1] "a" "c"
> s[c(TRUE, FALSE, TRUE, FALSE, FALSE)]
[1] "a" "c"
> a<-1:5
> a<3
[1] TRUE TRUE FALSE FALSE FALSE
> s[a<3]
[1] "a" "b"
> s[a>1 & a<3]
[1] "b"
> s[a==2]
[1] "b"
```

Operators

- arithmetic
+, -, *, /, ^
 - comparison
<, >, <=, >=, ==, !=
 - logical
!, &, |, xor
- (equal to, not equal to)
- these always return logical values !
(TRUE, FALSE)

Exercise

- Create the patients data frame using the instructions in the slides.
- Add three new variables to your data frame: country, continent, and height. Make up the data. Make country a character vector but continent a factor.
- Try the **summary** function on your data frame. What does it do? How does it treat vectors (numeric, character, logical) and factors? (What does it do for matrices?)
- Use logical indexing to select the following patients:
 - Patients under 40
 - Patients who give consent to share their data
 - Men who weigh as much or more than the average European male (70.8 kg)

Logical indexing answers

- Patients under 40:

```
> patients[patients$Age<40,]
```
- Patients who give consent to share their data:

```
> patients[patients$Consent==TRUE,]
```
- Men who weigh as much or more than the average European male (70.8 kg):

```
> patients[patients$Sex=="Male" & patients$Weight<=70.8,]
```
