

A practical tutorial on S4 programming

LAURENT GATTO*

Cambridge Center for Proteomics
University of Cambridge

June 21, 2013

Contents

1	Introduction	2
2	The microarray example	2
3	Using OO programming	6
4	The MArray class	7
5	MArray methods	9
5.1	The show method	9
5.2	Accessors	11
5.3	The sub-setting operation	12
5.4	The validity method	13
5.5	A replace method	13
6	Introspection	13
7	Conclusion	13

*lg390@cam.ac.uk

1 Introduction

This document¹ introduces the R object-oriented programming paradigm using the microarray as a use case. The introduction is purely practical and does not aim for an exhaustive guide to R object-oriented programming. We will concentrate on the S4 system and only mention the older S3 system and the recent S4 reference class infrastructure. See the appropriate literature, `?ReferenceClasses` or our more thorough introduction to OO programming² and references therein for more details.

In section 2, we present a solution on how to represent microarray data in R using simple data types and conclude with some issues with this implementation. In section 3, we introduce fundamental concepts of OO programming and introduce how OO programming is implemented in S4 (and S3) system.

2 The microarray example

We assume the reader is familiar with the concept of microarrays and the type of data that is obtained from such an experiment. Before embarking in any serious programming task, in particular when modelling data and defining data structures (using a OO class or not), to carefully think about how to best represent and store the data.

Exercise 1: Based on your understanding of microarrays and the kind of data that is to be used to computational analysis, think of what is going to be needed to describe an experiment and what the types of data structures available in R (`data.frame`, `matrix`, `vector`, ...) are most appropriate. Ideally, one would want everything (data, meta-data, ...) to be stored together as a single variables.

There are of course multiple valid solutions to the above question. Below are three pieces of information that consider essential along with their respective R data structure.

- We choose to represent the microarray results as a `matrix` of size $n \times m$, where n is the number of probes on the microarray and m is the number of samples. The matrix that stores the intensities (these could also represent fold-changes) is named `marray`.
- The sample annotation (meta-data) is described using a `data.frame` with exactly m rows and any number of columns. It is named `pmeta`.

¹The latest version of this document is available on its github repository <https://github.com/lgatto/S4-tutorial>.

²<https://github.com/lgatto/roo>

- The feature (probe) annotation (meta-data) is described using a `data.frame` with exactly n rows and any number of columns. Let's call it `fmeta`.

```
> n <- 10
> m <- 6
> marray <- matrix(rnorm(n * m, 10, 5), ncol = m)
> pmeta <- data.frame(sampleId = 1:m,
+                      condition = rep(c("WT", "MUT"), each = 3))
```

We will manually use the same names for intensity matrix columns and the sample meta-data rows as well as the matrix rows and feature meta-data row. Finally, to keep these pieces of information together, they will all be combined into a `list` that will represent our microarray experiment.

```
> rownames(pmeta) <- colnames(marray) <- LETTERS[1:m]
> fmeta <- data.frame(geneId = 1:n,
+                    pathway = sample(LETTERS, n, replace = TRUE))
> rownames(fmeta) <-
+   rownames(marray) <- paste0("probe", 1:n)
> maexp <- list(marray = marray,
+              fmeta = fmeta,
+              pmeta = pmeta)
> rm(marray, fmeta, pmeta)
> str(maexp)
```

List of 3

```
$ marray: num [1:10, 1:6] 6.87 10.92 5.82 17.98 11.65 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:10] "probe1" "probe2" "probe3" "probe4" ...
.. ..$ : chr [1:6] "A" "B" "C" "D" ...
$ fmeta : 'data.frame': 10 obs. of 2 variables:
..$ geneId : int [1:10] 1 2 3 4 5 6 7 8 9 10
..$ pathway: Factor w/ 8 levels "E","F","L","M",...: 8 4 4 1 7 3 5 2 2 6
$ pmeta : 'data.frame': 6 obs. of 2 variables:
..$ sampleId : int [1:6] 1 2 3 4 5 6
..$ condition: Factor w/ 2 levels "MUT","WT": 2 2 2 1 1 1
```

We can access the respective elements of our microarray experiment with the `$` operator.

```
> maexp$pmeta
```

	sampleId	condition
A	1	WT
B	2	WT
C	3	WT
D	4	MUT
E	5	MUT
F	6	MUT

```
> summary(maexp$marray[, "A"])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
5.82	7.27	11.30	10.70	12.80	18.00

```
> wt <- maexp$pmeta[, "condition"] == "WT"
> maexp$marray["probe8", wt]
```

A	B	C
13.692	14.719	2.646

```
> maexp[["marray"]][["probe3", !wt]
```

D	E	F
11.94	13.48	11.71

Exercise 2: But what if we want to subset the experiment. How would we extract the 10 first probes for the 3 first samples?

We have to manually subset the individual elements of our list, making sure that the number of rows of the `marray` and `fmeta` elements remain identical as well as the number of columns of `marray` and the number of columns of `pmeta`.

```
> x <- 1:5
> y <- 1:3
> marray2 <- maexp$marray[x, y]
> fmeta2 <- maexp$fmeta[x, ]
> pmeta2 <- maexp$pmeta[y, ]
> maexp2 <- list(marray = marray2, fmeta = fmeta2, pmeta = pmeta2)
> rm(marray2, fmeta2, pmeta2)
> str(maexp2)
```

```
> boxplot(maexp$marray)
```




Figure 1: Boxplot representing the intensity distributions of the 10 probes for the 6 samples.

```
List of 3
```

```
$ marray: num [1:5, 1:3] 6.87 10.92 5.82 17.98 11.65 ...
..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:5] "probe1" "probe2" "probe3" "probe4" ...
.. ..$ : chr [1:3] "A" "B" "C"
$ fmeta : 'data.frame': 5 obs. of 2 variables:
..$ geneId : int [1:5] 1 2 3 4 5
..$ pathway: Factor w/ 8 levels "E","F","L","M",...: 8 4 4 1 7
$ pmeta : 'data.frame': 3 obs. of 2 variables:
..$ sampleId : int [1:3] 1 2 3
```

```
..$ condition: Factor w/ 2 levels "MUT","WT": 2 2 2
```

The above solution does not provide a clean syntax. As a user, we have to know the names or positions of the respective elements of the microarray list elements to directly access the parts of interest. Finally, a simple operation like sub-setting the microarray experiment is very cumbersome and prone to errors.

3 Using OO programming

Object-oriented programming is based on two important concepts, abstraction and encapsulation. We want to represent the microarray concept in a way that makes most sense to the users without distracting them with unnecessary technicalities. These technicalities refer to the underlying implementation. Do the users really need to know that we used a list and that the first element, called `marray` is the matrix? We want the users to comprehend microarrays in R like they know them in real life, i.e. manipulate the abstract concept microarray while keeping all the underlying technical details, the implementation, hidden, or encapsulated.

These goals are achieved in two steps. First, we defined a class that represents (abstracts) the concept of a microarray. This is very similar to what we have done with the `list` above (the S3 system does use list), but we will use a more elaborated approach that, although more verbose, provides numerous benefits that will be described in the next sections. The class represents a data container and is defined on its own. An instance of a specific class, that contains data arranged in the specific container, is called an object.

Once we have created a class, we will want to defined a set of specific behaviours, that make sense in the eyes of the users. These behaviours will be implemented by special functions, called methods. Methods are functions that tune their behaviour based on the class of their input. You have already observed this in your every day usage of R : whether we ask to produce the boxplot of a `matrix` (for example `boxplot(maexp[[1]])`) or provide a `data.frame` and a `formula` like `boxplot(sampleId ~ condition, data = maexp[[3]])`, R automatically does the right thing.

It now becomes obvious that we have two different kind of roles. The *developer* is the one that creates the class and knows the implementation and the *user* is the one that uses the class without knowing, or needing to know, its actual underlying representation.

4 The MArray class

We can define a class with the `setClass` function. Our class is defined by a name, `MArray`, and a content. The different elements of an S4 class are called slots³.

```
> MArray <- setClass("MArray",  
+                   slots = c(marray = "matrix",  
+                             fmeta = "data.frame",  
+                             pmeta = "data.frame"))
```

The `setClass` function returns a special function called a constructor, that can be used to create an instance of the class.

```
> ## an empty object  
> MArray()  
  
An object of class "MArray"  
Slot "marray":  
<0 x 0 matrix>  
  
Slot "fmeta":  
data frame with 0 columns and 0 rows  
  
Slot "pmeta":  
data frame with 0 columns and 0 rows  
  
> ma <- MArray(marray = maexp[[1]],  
+             pmeta = maexp[["pmeta"]],  
+             fmeta = maexp[["fmeta"]])  
> class(ma)  
  
[1] "MArray"  
attr("package")  
[1] ".GlobalEnv"  
  
> ma  
  
An object of class "MArray"  
Slot "marray":
```

³Note that the usage of `slots` to define the representation of the class is the preferred way to define a class; the `representation` function is deprecated from version 3.0.0 and should be avoided.

	A	B	C	D	E	F
probe1	6.868	17.559	14.59489	16.793	9.177	11.991
probe2	10.918	11.949	13.91068	9.486	8.733	6.940
probe3	5.822	6.894	10.37282	11.938	13.485	11.706
probe4	17.976	-1.073	0.05324	9.731	12.783	4.353
probe5	11.648	15.625	13.09913	3.115	6.556	17.165
probe6	5.898	9.775	9.71936	7.925	6.463	19.902
probe7	12.437	9.919	9.22102	8.029	11.823	8.164
probe8	13.692	14.719	2.64624	9.703	13.843	4.779
probe9	12.879	14.106	7.60925	15.500	9.438	12.849
probe10	8.473	12.970	12.08971	13.816	14.406	9.325

Slot "fmeta":

	geneId	pathway
probe1	1	Z
probe2	2	M
probe3	3	M
probe4	4	E
probe5	5	T
probe6	6	L
probe7	7	N
probe8	8	F
probe9	9	F
probe10	10	P

Slot "pmeta":

	sampleId	condition
A	1	WT
B	2	WT
C	3	WT
D	4	MUT
E	5	MUT
F	6	MUT

To access individual slots, we need to use the `@`. This is equivalent of using the `$` for a list.

```
> ma@pmeta
```

	sampleId	condition
A	1	WT
B	2	WT
C	3	WT
D	4	MUT
E	5	MUT
F	6	MUT

A	1	WT
B	2	WT
C	3	WT
D	4	MUT
E	5	MUT
F	6	MUT

But this is something we do not want a user to do. To access a slot like this, one needs to know its name, i.e. the underlying plumbing of the class. This breaks the notion of encapsulation. Instead, the developer will provide the user with specific accessor methods (see section 5.2) to extract (or update using a replace method, section 5.5) specific slots.

5 MArray methods

Before proceeding, we need to explain the concept of generic function. A generic function, or generic for short, is a function that dispatches methods to the appropriate class-specific implementation. A method `do` will implement behaviour for a specific class `A`, while another implementation of `do`, will define another behaviour for class `B`. The generic `do` is the link between the class and its dedicated implementation. If we have `do(a)` (where `a` is of class `A`), then the generic will make sure that the `A`-specific code of `do` will be executed.

Before we define a method with `setMethod`, we will always want to first check if such a method does not exist (in which case there is already a generic function), as illustrated with the `show` method in section 5.1. If it is the case, we write our new methods. If not, we first create the generic with `setGeneric` and then proceed with the method.

5.1 The `show` method

The `show` method (it is a method, as it exhibits custom behaviour depending on the class of its argument) is a very helpful one. It allows to define custom summary view of an object when we type its name in the console, instead of having all its (possibly very long content) displayed.

```
> show

standardGeneric for "show" defined from package "methods"

function (object)
```

```

standardGeneric("show")
<bytecode: 0x1807a98>
<environment: 0x17ea328>
Methods may be defined for arguments: object
Use showMethods("show") for currently available ones.
(This generic function excludes non-simple inheritance; see ?setIs)

> isGeneric("show")

[1] TRUE

> hasMethod("show")

[1] TRUE

```

As there is already a `show` generic function, we can immediately proceed with the method definition using the `setMethod` function. To do so we need a few things. First, we need to know for what class we implement the specific `show` method; this is the `MArray` class and will be passed as the `signature` argument in `setMethod`. We also need to know the argument names that are defined in the generic. These must match exactly, as we write a method for that specific generic. The arguments can be found by just typing the name of the generic (as in the previous) code chunk, look at its documentation or directly ask for the arguments with `args(show)`. We see that there is only one argument, `object` (naming the first argument of a generic `object` is a widely applied convention). This is the same name that we will have to use when writing the definition of our method.

```

> setMethod("show",
+           signature = "MArray",
+           definition = function(object) {
+               cat("An object of class ", class(object), "\n", sep = "")
+               cat(" ", nrow(object@marray), " features by ",
+                   ncol(object@marray), " samples.\n", sep = "")
+               invisible(NULL)
+           })

[1] "show"

> ma

An object of class MArray
 10 features by 6 samples.

```

5.2 Accessors

As mentioned above, we want to provide customised and controlled access to the class slots. This does not prevent us, as developers, to use the `@` accessor, but does not force others to know the implementation details.

Let's create an accessor for the `marray` slot and call the accessor `marray`. There is not harm in naming the slot and its accessor but there is no constraint in doing so. There is no such method or generic; just typing `marray` will tell you that no such object is found. Below, we create a new generic function with `setGeneric`. We define the name of our new generic as well as the name of the argument(s) that will have to be re-used when defining our class-specific method.

```
> setGeneric("marray", function(object) standardGeneric("marray"))  
  
[1] "marray"
```

We now proceed in the same way as previously, using `setMethod`. The definition of our method (i.e. the actual code that will be executed) is very basic and of course uses the `@`.

```
> setMethod("marray", "MArray",  
+          function(object) object@marray)  
  
[1] "marray"  
  
> marray(ma)
```

	A	B	C	D	E	F
probe1	6.868	17.559	14.59489	16.793	9.177	11.991
probe2	10.918	11.949	13.91068	9.486	8.733	6.940
probe3	5.822	6.894	10.37282	11.938	13.485	11.706
probe4	17.976	-1.073	0.05324	9.731	12.783	4.353
probe5	11.648	15.625	13.09913	3.115	6.556	17.165
probe6	5.898	9.775	9.71936	7.925	6.463	19.902
probe7	12.437	9.919	9.22102	8.029	11.823	8.164
probe8	13.692	14.719	2.64624	9.703	13.843	4.779
probe9	12.879	14.106	7.60925	15.500	9.438	12.849
probe10	8.473	12.970	12.08971	13.816	14.406	9.325

If we change the underlying implementation by changing the name of the slot or using an environment instead of a matrix, the `ma@marray` is going to break.

However, when providing accessors, we can echo the changes in the accessor implementation without affecting the users' behaviour or existing scripts.

5.3 The sub-setting operation

Let's now encapsulate the sub-setting of an `MArray` object in a proper method so facilitate this simple operation. In `R`, the default subsetting operator is `[`. Although its syntax looks like it is special, the operator is just a normal function with a bit of extra syntactic sugar.

```
> letters[1:3]

[1] "a" "b" "c"

> letters[1:3]

[1] "a" "b" "c"
```

If you type `'['` in your `R` console, you will see that this is a primitive function. These internally implemented functions have a special property that, although not explicitly generic functions, they get automatically promoted to generics when a method of the same name is defined. In other words, must not create a generic (this would break `[` in all the other cases) and can directly proceed with the method definition.

We can thus implement a specific behaviour of `[` for the `MArray` class. The documentation `help("[")` shows that, in addition to `x`, the object to be subset, we also have to take the `i` and `j` indices into account and the `drop` argument. When an argument is not relevant, we specify this by declaring that it is `"missing"`.

```
> setMethod("[", "MArray", function(x, i, j, drop = "missing") {
+   .marray <- x@marray[i, j]
+   .pmeta <- x@pmeta[j, ]
+   .fmeta <- x@fmeta[i, ]
+   MArray(marray = .marray, fmeta = .fmeta, pmeta = .pmeta)
+ })

[1] "["

> ma[1:5, 1:3]

An object of class MArray
5 features by 3 samples.
```

5.4 The validity method

5.5 A replace method

6 Introspection

```
> slotNames(ma)

[1] "marray" "fmeta"  "pmeta"

> getClass("MArray")

Class "MArray" [in ".GlobalEnv"]

Slots:

Name:      marray      fmeta      pmeta
Class:     matrix data.frame data.frame
```

Finding stuff with `showMethods`, `getMethod()`.

7 Conclusion

A little overhead in the beginning for improved stability and usability in the future. This however implied to think about the needs of the users role before you write code.

Session information

All software and respective versions used to produce this document are listed below.

- R Under development (unstable) (2013-06-16 r62969),
x86_64-unknown-linux-gnu
- Locale: LC_CTYPE=en_GB.UTF-8, LC_NUMERIC=C, LC_TIME=en_GB.UTF-8,
LC_COLLATE=en_GB.UTF-8, LC_MONETARY=en_GB.UTF-8,
LC_MESSAGES=en_GB.UTF-8, LC_PAPER=C, LC_NAME=C, LC_ADDRESS=C,
LC_TELEPHONE=C, LC_MEASUREMENT=en_GB.UTF-8, LC_IDENTIFICATION=C
- Base packages: base, datasets, graphics, grDevices, methods, stats, utils
- Other packages: codetools 0.2-8, knitr 1.2

- Loaded via a namespace (and not attached): digest 0.6.3, evaluate 0.4.3, formatR 0.7, stringr 0.6.2, tools 3.1.0